

Root Server Selection in Recursive DNS Software

Through dynamic debugging, we review the root query policies in four types of popular recursive DNS software: BIND 9 [3], Unbound [5], Knot Resolver [1] and PowerDNS Recursor [6]. All software reuses *authoritative server (NS) selection algorithms* for root queries, and in the case of the DNS root, root servers are their NSes. Below we list pseudo-code and detailed analysis of each type of DNS software.

1 BIND 9

Algorithm 1 shows the selection algorithm of BIND 9. At startup, BIND launches priming queries [4] to find the information of all root servers. For a set of NSes BIND maintains their *adjusted RTTs*, which are dynamically updated using the Exponentially Weighted Moving Average (EWMA) [2], taking the response status (e.g., error and timeout rate), latency and lameness of each NS as factors. From this set BIND will select and query the NS with the **lowest adjusted RTT**. NSes that have not been queried will be given a default adjusted RTT value ranging from 1 to 32 microseconds, such that in early stages they will quickly be tried and have their RTTs updated. After each round BIND also decreases the adjusted RTTs for NSes that are not selected, giving them chances to be selected in less cases. IPv6 addresses are preferred to IPv4.

2 Knot Resolver

Algorithm 2 shows the selection algorithm of Knot Resolver. Knot Resolver launches priming queries at startup. It uses the **ϵ -Greedy algorithm** to select from a list of NSes, giving priorities to NSes that support IPv6, are never tried, with less query failures and lower delay. In around 1/20 cases, Knot Resolver will randomly select one from all untried NSes.

3 Unbound

Algorithm 3 shows the selection algorithm of Unbound. Unlike other software, Unbound does not launch prime queries at startup and only queries root servers when necessary. An adjusted RTT is also maintained for each NS, and NSes with bad status (e.g., lame DNSSEC, long latency or timeout) will receive penalty. Each untried NS is given an initial adjusted RTT of 376ms. Unbound **randomly**

selects from a list of valid NSes and removes an NS from consideration only when its adjusted RTT is over 400ms longer than the least-latent. However, in the case of DNS root, only few locations across the globe witness a delay of over 400ms thus Unbound will tend to select root servers randomly.

Note that the algorithm described here is enabled by default. Through custom configuration, Unbound can also select the best NS from a valid list at a given probability.

4 PowerDNS Recursor

Algorithm 4 shows the selection algorithm of PowerDNS Recursor. PowerDNS Recursor launches prime queries at startup. Similiar to BIND, PowerDNS Recursor maintains the adjusted RTT for each NS and dynamically update them using the EWMA. The default adjusted RTT value is 0ms, thus in early stages PowerDNS Recursor selects NSes randomly. If a query to an NS fails, the server will be set as throttled and removed from consideration temporarily. **It selects NS with the lowest adjust RTT and decays the adjusted RTT for all servers with the same factor. The longer the query interval, the lower the decay factor and adjusted RTTs.** It also removes the RTT record of a NS periodically to ensure every record is up-to-date. Because root servers are typically not frequently queried, the interval will be long, resulting in low decay factor and adjusted RTTs for servers not selected in the current round.

References

1. CZ.NIC: Knot resolver (2021), <https://www.knot-resolver.cz/>
2. Hunter, J.S.: The exponentially weighted moving average. *Journal of quality technology* **18**(4), 203–210 (1986)
3. ISC: Bind 9 (2021), <https://www.isc.org/bind/>
4. Koch, P., Larson, M., Hoffman, P.E.: Initializing a DNS Resolver with Priming Queries. RFC 8109 (Mar 2017). <https://doi.org/10.17487/RFC8109>, <https://rfc-editor.org/rfc/rfc8109.txt>
5. Labs, N.: Unbound (2021), <https://www.nlnetlabs.nl/projects/unbound/about/>
6. PowerDNS: Powerdns recursor (2020), <https://www.powerdns.com/recursor.html>

Algorithm 1 BIND 9: Select the best NS candidate

```

1: function SORTALLNSES( $L_{NS}$ )
2:   for each  $ns \in L_{NS}$  do
3:     if  $ns$  is an IPv4 address then
4:        $tmp\_srtt$  of  $ns \leftarrow$   $\_srtt$  of  $ns$  + penalty value
5:     end if
6:   end for
7:   bubble sort  $L_{NS}$  by  $tmp\_srtt$  ▷ ascending
8:   return  $L_{NS}$ 
9: end function
10:
11: function FINDTHEBESTNS( $L_{NS}$ )
12:   for each  $ns \in L_{NS}$  do
13:     if  $ns$  isn't tried then
14:       set  $ns$  is tried
15:       return  $ns$ 
16:     end if
17:   end for
18:   return no more NS
19: end function
20:
21: function AFTERQUERY( $ns, rtt, L_{NS}$ )
22:   for each  $ns \in L_{NS}$  do
23:     if  $ns$  isn't tried then
24:        $\_srtt$  of  $ns \leftarrow (\_srtt$  of  $ns \cdot 2^9 - \_srtt$  of  $ns)/2^9$ 
25:     end if
26:   end for
27:   if query succeeded then
28:      $v \leftarrow rtt$ 
29:   else
30:      $v \leftarrow$  a penalty value
31:   end if
32:    $\_srtt$  of  $ns \leftarrow \_srtt$  of  $ns \cdot (1 - a) + a \cdot rtt$  ▷  $a$  is depended on status of  $ns$ 
33: end function

```

Algorithm 2 Knot Resolver: Select the best NS candidate

```

1: function FINDHIGHERPRIORITYNS(Two NS candidates)
2:   if one NS has IPv6 address, the other one don't and IPv6 network enabled
   then
3:      $ns \leftarrow$  the NS has IPv6 address
4:   else if one NS is never tried before, the other one has tried then
5:      $ns \leftarrow$  the NS which is never tried before
6:   else if one NS has less error in previous probes then
7:      $ns \leftarrow$  the NS has less error
8:   else
9:     if one NS has lower  $\_srtt$  then
10:       $ns \leftarrow$  the NS has lower  $\_srtt$ 
11:     else
12:       return equal priority
13:     end if
14:   end if
15:   return  $ns$ 
16: end function
17:
18: function FINDTHEBESTNS( $L_{NS}$ )
19:   shuffle  $L_{NS}$  randomly
20:   quick sort  $L_{NS}$  by FINDHIGHERPRIORITYNS ▷ descending
21:    $r \leftarrow$  a random value,  $r \in [0, 1]$  ▷  $\epsilon$ -Greedy Selection
22:   if  $r < \frac{1}{20}$  then ▷ Explore
23:      $ns \leftarrow$  a random untried NS in  $L_{NS}$ 
24:   else ▷ Exploit
25:      $ns \leftarrow L_{NS}[0]$ 
26:   end if
27:   if  $ns$  has errors before and still has untried NS in  $L_{NS}$  then
28:      $ns \leftarrow$  a random untried NS in  $L_{NS}$ 
29:   end if
30:   return  $ns$ 
31: end function
32:
33: function AFTERQUERY( $ns, rtt$ )
34:   if  $ns$  is never tried before then
35:      $\_srtt$  of  $ns \leftarrow 400\text{ms}$ 
36:     set  $ns$  has tried
37:   else
38:     if query succeeded then
39:       use  $rtt$  to update  $\_srtt$  of  $ns$  according to section 3 in RFC6298
40:     else
41:        $error$  of  $ns \leftarrow error$  of  $ns + 1$ 
42:     end if
43:   end if
44: end function

```

Algorithm 3 Unbound: Select the best NS candidate

```

1: function FINDTHEBESTNS( $L_{NS}$ )
2:   for each  $ns \in L_{NS}$  do
3:     if  $ns$  is bogus or lame or in unsupported network or not allowed to be
       queried then
4:       remove  $ns$  from  $L_{NS}$ 
5:       continue
6:     else if  $ns$  is never tried before then
7:        $tmp\_srtt$  of  $ns \leftarrow 376ms$ 
8:     else if  $ns$  is in a bad status (e.g., dnssec lame, huge timeout) then
9:        $tmp\_srtt$  of  $ns \leftarrow$  a corresponding penalty value
10:    else
11:       $tmp\_srtt$  of  $ns \leftarrow \_srtt$  of  $ns$ 
12:    end if
13:     $best\_srtt \leftarrow \text{MIN}(best\_srtt, tmp\_srtt \text{ of } ns)$ 
14:  end for
15:  for each  $ns \in L_{NS}$  do
16:    if  $tmp\_srtt$  of  $ns > best\_srtt + 400ms$  then
17:      remove  $ns$  from  $L_{NS}$ 
18:    end if
19:  end for
20:  return a random NS in  $L_{NS}$ 
21: end function
22:
23: function AFTERQUERY( $ns, rtt$ )
24:  set  $ns$  is tried in this turn of query
25:  if query succeeded then
26:    use  $rtt$  to update  $\_srtt$  of  $ns$  according to section 3 in RFC6298
27:  else
28:    record corresponding status(as described in FINDTHEBESTNS) of  $ns$ 
29:  end if
30: end function

```

Algorithm 4 PowerDNS: Select the best NS candidate

```

1:  $now \leftarrow$  current time (seconds)
2:
3: function SORTALLNSES( $L_{NS}$ )
4:   for each  $ns \in L_{NS}$  do
5:     if  $ns$  isn't set throttled then
6:        $d \leftarrow last$  of  $ns - now$ 
7:        $\_srtt$  of  $ns \leftarrow \_srtt$  of  $ns \cdot \exp(d/60)$ 
8:        $last$  of  $ns \leftarrow now$ 
9:     end if
10:  end for
11:  shuffle  $L_{NS}$  randomly
12:  stable sort  $L_{NS}$  by  $\_srtt$   $\triangleright$  ascending order
13:  return  $L_{NS}$ 
14: end function
15:
16: function FINDTHEBESTNS( $L_{NS}$ )
17:  for each  $ns \in L_{NS}$  do
18:    if  $ns$  isn't tried then
19:      return  $ns$ 
20:    end if
21:  end for
22:  return no more NS
23: end function
24:
25: function AFTERQUERY( $ns, rtt$ )
26:  set  $ns$  is tried in this turn of query
27:  if query succeeded then
28:    if  $ns$  doesn't have  $\_srtt$  record then
29:       $\_srtt$  of  $ns \leftarrow rtt$ 
30:    else
31:       $d \leftarrow last$  of  $ns - now$ 
32:       $a \leftarrow \exp(d)/2$ 
33:       $\_srtt$  of  $ns \leftarrow \_srtt$  of  $ns \cdot a + (1 - a) \cdot rtt$ 
34:       $last$  of  $ns \leftarrow now$ 
35:    end if
36:  else
37:    set  $ns$  is throttled.
38:  end if
39: end function
40:
41: function HOUSEKEEPING( $L_{NS}$ )  $\triangleright$  An independent thread to remove the status of
    NSes periodically
42:  for each  $ns \in L_{NS}$  do
43:    every 5 seconds, remove the throttled status of  $ns$ 
44:    every 200 seconds, remove  $\_srtt$  of  $ns$ , whose  $now - last > 300$ 
45:  end for
46: end function
  
```
